

## Intelligent Agents

Rationality: achieving maximum utility by some pre-defined metric or set of goals/intentions.  
depends on the usefulness of the choice and not on the process that led to that choice  
e.g., rational process for playing tic-tac-toe by tabulating game states  
has no decision process, but acts rationally

An agent has sensors which take in percepts and actuators which effect actions  
a percept is a set of perceptual inputs at a fixed point in time  
a percept sequence is composed of percepts

An agent function is  $f : P^* \rightarrow A$  where  $P^*$  is the percept sequence and  $A$  the set of actions  
the agent program is the specific architecture which implements this function  
not all agent functions can be implemented by some agent program  
e.g. halting problems, NP-hard problems, “too-large” problems (e.g. chess)

Performance measure: an objective criterion for the success of behavior of an agent

Rational agent: maximizes expected performance measure by its actions  
given the prior knowledge and percept sequence available to it  
e.g. vacuum world: 2 squares, could be dirty or clean  
action function: suck if dirty, move if clean  
under the measure of most clean squares/time period, this is rational  
if we seek to minimize movements as well, this is irrational

Autonomy: the ability to function beyond the prior knowledge of the designer

The task environment: performance measure, environment, actuators, and sensors  
partially vs. fully observable (perceive all aspects relevant to choice of action)  
stochastic vs. deterministic (next state a function only of current state, action)  
strategic: deterministic but for the actions of other agents  
episodic vs. sequential (current decision could affect future decisions)  
static vs. dynamic (environment can change while agent deliberates)  
semidynamic: environment doesn't change with time but performance score does  
semidynamic e.g. chess with a clock  
discrete vs. continuous (can apply to state, time, percepts/actions)  
single vs. multiagent: e.g. competitive multiagent, cooperative multiagent

Agent structure

the agent program only takes in the current percept  
the agent function maps from the entire percept history

A simple reflex agent uses only the current percept in its decision process  
responds via condition-action rules  
requires full observability to work effectively

A state-based reflex agent maintains internal state using the percept history

needs what is called a model of the environment  
knowledge about how it evolves, how actions affect it  
also called a model-based reflex agent

A goal-based agent seeks to achieve states which are considered favorable

Unlike reflex agents, makes use of foresight in its decision process

A utility-based agent evaluates future states using a utility function

Seeks to maximize utility of future states

A learning agent can be broken down roughly into four components

A learning element is responsible for improving the agent

A critic evaluates current performance and informs the learning element  
(uses a fixed performance standard)

A performance element selects actions

A problem generator suggests exploratory actions

Two of these are reflexive, the next two are planning-based.

Planning agents predict consequences of actions using a transition model.

Agents can vary in the degree to which they deliberate

one extreme: carefully construct a complex plan

another: start with a simple plan and rapidly correct as complications arise

## Search Problems

We consider problem-solving agents, which set out to achieve some desirable state

an uninformed search algorithm has no idea of where to look for solutions

relies only upon the problem definition

An agent plan in a search problem will first formulate, then search, then execute

Problem formulation:

simplifies environment, abstracts away unnecessary information

helps organize the behavior of the agent

Components of a well-defined problem:

(1) Initial state

(2) Possible actions: i.e., a successor function  $f : \{\text{states}\} \rightarrow \{(\text{action}, \text{consequent state})\}$

These first two implicitly define a state space, and a state space graph

State space graph has states as nodes and actions as edges

(3) Goal test

(4) Path cost: (e.g. induced by edge costs of a state space graph)

A path is a sequence of actions.

A solution is a path from an initial state to a goal state.

An optimal solution has lowest path cost.

Environments we are considering are static, observable, discrete, and deterministic.

Generalized tree search

envision the search space as a tree (states can be revisited)

nodes have: state, parent-node, action leading to that node, path-cost, and depth

the root of the search tree corresponds to the initial state of the problem

uses a fringe, initialized to contain only the root node  
the fringe contains nodes which have been generated but not yet expanded

At each node in the process:

perform goal test on the node in question, if succeeds, return corresponding solution  
expand node by i.e. applying successor function to generate new nodes  
if no such nodes exists, declare failure  
choose which node to analyze next according to some search strategy

Evaluation of a search strategy.

Completeness: whether or not it finds a solution, if a solution exists.

Optimality: if it finds a solution, whether or not such a solution is of lowest cost.

Complexity: space (nodes stored at any given time) or time (total nodes explored)

Factors often used to evaluate the complexity of a given algorithm.

Branching factor  $b$ : the maximum number of successors of any node.

Depth  $d$ : the minimum depth among goal nodes.

Maximum path length  $m$ : can be infinite.

Types of uninformed (aka blind) search strategies

contrast: informed search strategies take advantage of heuristics

cannot solve searches of exponential complexity for all but the smallest problems

Breadth-first search uses a queue (FIFO) as its fringe.

Satisfies completeness (if branching factor finite) but not optimality.

Identical space and time complexity (holds all nodes in state until goal is found)

Exponential complexity:  $\mathcal{O}(b^{d+1})$

In practice, spacial complexity too hard: very difficult to accrue enough memory

Uniform-cost search orders its fringe by path cost.

As long as each step cost  $> \epsilon > 0$ , satisfies completeness *and* optimality.

Again, same time and space complexity

Given  $\epsilon$  minimum action cost,  $C^*$  optimal solution cost, complexity is  $\mathcal{O}(b^{\lceil C^*/\epsilon \rceil})$

Depth-first search uses a stack (LIFO) as its fringe

Node storage is not exponential, space complexity is  $\mathcal{O}(bm)$ .

Backtracking search: nodes store set of successors, search expands only one at a time.

Backtracking search can reduce the space complexity to  $\mathcal{O}(m)$ .

Worst-case time complexity  $\mathcal{O}(b^m)$  where potentially  $m \gg d$  and even  $m = \infty$ .

Complete if  $m < \infty$ , but not optimal.

Depth-limited search

specifies a depth limit  $l$ , and performs a depth-first search up to that limit

if  $l < d$ , will be incomplete, and if  $d < l$ , can be non-optimal

has a time complexity of  $\mathcal{O}(b^l)$  and a space complexity of  $\mathcal{O}(bl)$ .

the diameter of the state space: min number of actions between any two states

is a great depth limit, but, we don't necessarily know the diameter, a priori

Depth-first search by iterative deepening

apply depth-limited search with  $l$  ranging over  $\mathbb{N}$

like DFS, has good memory (spatial complexity) at  $\mathcal{O}(bd)$

like BFS, if  $b$  is finite will be complete, and if path cost corresponds to depth, optimal

repeated state generation not costly: most nodes in the bottom level  
generated nodes:  $d(b) + (d-1)b^2 + \dots + (1)b^d$  (factor of  $d$  is node repetition)  
hence the time complexity is  $\mathcal{O}(b^d)$

better than BFS ( $\mathcal{O}(b^{d+1})$ ): IDS generates no nodes beyond the solution depth

When search space is large and solution depth unknown, depth-first IDS is generally best.

Bidirectional search run simultaneous searches from initial state and goals

Concludes when the searches meet.

Time and complexity  $\mathcal{O}(b^{d/2})$ , is reduced.

However, requires effective computation of predecessors: often non-trivial/impossible.

Avoiding repetition of states: Graph Search

adds a closed list to tree search: a list of all nodes which have been expanded

the term open list is sometimes used to refer to the fringe

current nodes which match a node on the closed list are discarded

possibly suboptimal, if search methods can reach nodes at a non-optimal cost first

thus uniform cost graph search is optimal but IDS graph search may not be

the closed list increases space requirements, possibly to unfeasibility

behavior of closed list is such that the memory used is proportional to the runtime

Partial Information Search

Sensorless Problems

know consequences of actions and the possible states

can coerce world into a particular state, with some cleverness

uses a belief state: a set of states currently regarded as possible

a solution is a path to a belief state in which all its members satisfy the goal test

this approach can be analogously applied to nondeterministic problems

Contingency Problems

agent can obtain new information from sensors after performing actions

solution: use an unfixed action sequence, with dependencies on percepts received

the action at each node will depend on all percepts received up to that point

agent can act before finding a guaranteed plan

approach referred to as interleaving search and execution

does not need to account for *all* contingencies, simply the ones that occur

## Informed Search and Exploration

Informed search uses problem-specific knowledge beyond the definition of the problem

Approach: select nodes for expansion based on an estimate of distance to goal

Use a priority queue ordered by some evaluation function  $f$

Called best-first search: since we order by nodes which seem best

A heuristic  $h : \{\text{nodes}\} \rightarrow \mathbb{R}^+$  estimates cost of cheapest path to a goal node

Greedy best-first search uses  $f = h$

susceptible to false starts, dead ends

same defects as DFS: not optimal, incomplete, worst-case time/space complexity  $\mathcal{O}(b^m)$   
 $m$  the maximum depth of search space

A\* search uses  $f = g + h$  where  $g$  is the path cost to a node  
hence  $f$  is the estimated cost of cheapest solution through  $n$   
for tree search, if  $h$  is admissible (never overestimates), A\* will be optimal  
for graph search, need to ensure the first generated path is optimal  
this shall occur if  $h$  is additionally consistent/monotone  
consistent  $h$  satisfies, for nodes  $n, n'$  and action  $n \xrightarrow{a} n'$ ,  $h(n) \leq c(a) + h(n')$   
satisfies a triangle inequality; i.e.,  $h$  must be a metric on the state space  
A\* is optimally efficient for a given  $h$ :  $\nexists$  an optimal, more efficient algorithm  
however, nodes in goal contour search space still increase exponentially,  
unless  $|h - h^*| \leq \mathcal{O}(\log(h^*(n)))$  where  $h^*$  is the true cost  
generally  $|h - h^*| = \mathcal{O}(h^*(n))$  at best  
thus, it's often impractical to insist on finding an optimal solution  
A\* keeps all generated nodes in memory  $\rightarrow$  impractical for large-scale problems

Improvements on A\* with respect to space complexity

Can try iterative deepening A\* (IDA\*) using  $f = g + h$  rather than  $d$  as the cutoff  
but this incurs substantial overhead

Recursive best-first search

tracks the  $f$ -value of the current best alternative path  
winds back to this alternative path once  $f$  considered exceeds that stored  $f$   
optimal if  $h$  is admissible, with space complexity  $\mathcal{O}(bd)$   
time complexity can vary, depends on  $h$ , frequency of path changes  
can potentially explore a state multiple times (typical tree-search problem)  
stores only the value of  $f$  and  $\mathcal{O}(bd)$  nodes

MA\* (memory-bounded A\*) and SMA\* (simplified MA\*) make use of all memory  
description of SMA\*: keep expanding until memory is full  
once memory is full, replace the worst (by  $f$ ) node in memory, breaking ties by age  
SMA\* is complete if there is enough memory to hold the shortest path to a goal  
practically, probably best for a graph state space and non-uniform path costs  
if too much switching, problems that A\* would solve become intractable for SMA\*  
time  $\leftrightarrow$  space tradeoff

Learning to search better

metalevel learning algorithm analyzes current method, seeks improvements

Heuristics

**This is where we at.**

9/8

A heuristic  $h_a$  is dominant over  $h_c$ ,  $h_a \geq h_c$  if  $\forall n \ h_a(n) \geq h_c(n)$

A dominant heuristic of an admissible heuristic, if still admissible, is preferable

A maximum over admissible heuristics is a dominant heuristic, and still admissible